

The Autonomous Picking & Palletizing (APPLE) Robot: A Research Platform for Intralogistics Applications

Robert Krug^{*}, Todor Stoyanov^{*}, Vinicio Tincani[‡], Henrik Andreasson^{*},
Rafael Mosberger^{*}, Gualtiero Fantoni[‡], Achim J. Lilienthal^{*} and Antonio Bicchi[‡]

Abstract—In this work we present a research platform for fully autonomous commissioning tasks in intralogistics settings. The robot comprises a nonholonomic mobile base and a manipulation system consisting of a lightweight arm and an under-actuated gripper with active surfaces. The platform is capable of autonomously picking up pallets and loading them with unstructured goods in a manner which is safe for human workers sharing the environment. Target object handling is accomplished via a novel, redundant grasp representation which allows for redundancy in the gripper pose placement. This redundancy is exploited by an optimization-based control system which generates reactive manipulator motions on-the-fly without the delays occurring in sense-plan-act architectures.

I. INTRODUCTION

The increasing need for fast and flexible commissioning (*i.e.*, order picking and collection of unstructured goods from storage compartments in warehouses) in logistic scenarios has created substantial interest for autonomous robotic solutions. One of the main arguments for automating this task is that the dull and strenuous nature of commissioning could cause mental and physical illness in human workers. As a result, the determination within the logistics sector to invest in this area is high and substantial efforts are made for the humanization of workstations [1].

There exist partial solutions for the automated commissioning problem in controlled environments. The system described in [2] coordinates a fleet of Autonomous Ground Vehicles (AGVs) which transport shelves filled with goods to a human worker who picks the corresponding objects to complete the order. Key obstacles for a fully automatized solution applicable in general warehouse settings are the safe autonomous vehicle navigation in industrial environments co-populated by humans, as well as the autonomous grasping/manipulation of unstructured goods at reasonable cycle times.

In this work, we present the Autonomous Picking & Palletizing (APPLE) research platform (see Fig. 1) which we developed to address the following important sub-task chain which occurs during commissioning in prototypical warehouses: autonomous picking of goods from a storage location, subsequent placement on a standard EUR half-pallet and transport of the filled pallet to a target location. Furthermore, this process has to be carried out in a manner which is safe for humans operating in the same environment.

^{*} AASS Research Center; Örebro University; Studentgatan 1, 70182 Örebro, Sweden.

[‡] Interdepart. Research Center “E. Piaggio”; University of Pisa, Via Diotisalvi 2, 56100 Pisa, Italy.




Fig. 1. The APPLE platform: A KUKA LBR iiwa arm (3) is mounted on a retrofitted Linde CitiTruck AGV (7). For localization, a Velodyne HDL-32 Lidar (4) is used, human worker detection is carried out with the Reflex camera system (6). In order to detect pallets, we employ a Asus Xtion Pro Live structured light camera (5). The depicted grasping device (2) is a further developed and smaller version of the underactuated Velvet Fingers gripper described in [3]. Each of the grippers two fingers has a planar manipulator structure with two rotary joints and active surfaces which are implemented by conveyor belts on the inside of the two phalanges. These belts are used to assist in robust grasp acquisition as described in Section III-E. Object and pallet detection is done with a Structure IO device (1) which is mounted on the gripper's palm.

The platform's mobile base consists of a retrofitted nonholonomic Linde CitiTruck forklift AGV¹ which is able to detect and pick up pallets in designated loading zones. A KUKA LBR iiwa² lightweight arm which is fitted with an under-actuated gripper with conveyor belts on the inside of each finger is used for robust grasping and object manipulation. While this setup is not intended as a close-to-market solution for palletizing, the manipulator's force/torque sensing capabilities as well as the compliant and human-safe control possibilities allow to investigate various grasping and manipulation strategies.

¹<http://www.citi-truck.com>

²http://www.kuka-labs.com/en/service_robotics/lightweight_robotics/

In this paper, we outline our solution to the safe navigation of the APPLE platform in an industrial scenario co-populated by human workers wearing reflective vests. Furthermore, we introduce a novel grasp representation/planning scheme which is utilized for reactive, on-the-fly manipulator motion generation. This allows to exploit manipulator redundancies and offers several advantages to the commonly used sense-plan-act architectures. As a final contribution, we discuss our compliant grasp execution strategy which uses the active surfaces on the employed gripper to increase grasp robustness.

The remainder of this article is organized as follows: in Section II we outline the AGV navigation and motion planning scheme, as well as our solution for people detection in Section II-C. Section III is dedicated to the developed grasping and manipulation system with focus on object detection in Section III-B, grasp representation/planning in Section III-C and reactive manipulator motion generation in Section III-D. We then show early results and an evaluation of the APPLE system by means on a simplified commissioning task in Section IV before drawing conclusions in Section V.

II. AUTONOMOUS FORKLIFT

The mobile base is built upon a manual forklift which originally was equipped with motorized forks and a drive wheel only. The forklift has been retrofitted with a steering mechanism and a commercial AGV control system which is used to interface the original drive mechanism as well as the steering servo. To assure safe operation, the vehicle is equipped with a SICK S300 safety laser scanner³ and an industrial prototype system for detecting human workforce using reflective safety garments as described in Section II-C.

A. Challenges in Autonomous Navigation

The industry standard for autonomous navigation of forklifts is to use predefined trajectories which are either manually defined or learned through teaching-by-demonstration from a human operator [4], [5]. Although conceptually simple, fixing trajectories limits the pallet handling to occur only at predefined poses. Also, this only allows for overly simple strategies for handling unforeseen obstacles. The fundamental difficulties of motion planning for a forklift lie in the nonholonomic constraints, the large sweep area it needs to occupy (due to its very non-symmetrical footprint) while operating in a limited work space.

In order to obtain reliable localization in large dynamic warehouses with high accuracy, it is common to mount reflectors in the environment and to use a dedicated sensing device [6]. In this fashion, natural navigation has successfully been deployed in smaller entities where walls are commonly observed. However, without additional infrastructure, navigation in large and dynamic environments remains a challenge.

B. Navigation

The navigation module ensures that the forklift is capable of moving safely and autonomously through the work space environment to arbitrary load/unload poses with high accuracy. According to the AGV system provider Kollmorgen⁴, the required end pose accuracy for picking up pallets is ± 0.03 m in position and ± 1 degree in orientation. The navigation module comprises trajectory generation, tracking control and a localization system.

The online trajectory generation is done in two steps. First, a kinematically feasible path with discretized start and goal poses is generated using a lattice planner [7]. Second, this path is post-processed using a path smoother [8] which assures a smooth, collision-free and continuous trajectory. Subsequently, the obtained trajectory is tracked by a model predictive controller. For localization, a Velodyne HDL-32 3D laser scanner⁵ is utilized to construct a 3D map (using the 3D-NDT-OM map representation) of the static parts of the environment [9]. The map and odometry information is then used to localize the vehicle in the presence of dynamic entities using a dual timescale approach [10]. The complete navigation system has been implemented, extensively tested and successfully integrated on the APPLE demonstrator, a detailed description can be found in [8].

The current system for pallet detection and pickup requires a rough estimate of the location of the pallet (*i.e.*, a predefined pickup zone). In order to compute the final pose based on sensory data from an Asus Xtion Pro Live⁶ mounted on the AGV, a Signed Distance Function (SDF) tracker [11] is used with a predefined SDF model of the pallet. The tracking is done while driving towards the pickup zone and the trajectory is recomputed on the fly which, depending on the pose offset, may include a reverse driving operation.

C. People Detection

As the envisioned mobile manipulation system will operate in environments shared with human workers, people detection and human safety are important issues. In APPLE we address the problem by using the RefleX system we recently developed [12]. RefleX is a camera-based on-board safety system for industrial vehicles and machinery for detection of human workers wearing reflective vests worn as per safety regulations (see Fig. 2). The system was designed with industrial safety standards in mind and is currently being tested as an industrial prototype.

III. GRASPING & MANIPULATION SYSTEM

In this section we present an integrative approach for grasp representation/planning and motion generation. The main idea is to provide a functional representation of grasps as intervals in task space, which allows redundancy in the gripper pose prior to executing the grasp. Subsequently, we leverage the obtained redundancy to generate reactive

³<http://www.sick.com/>

⁴<http://www.kollmorgen.com/>

⁵<http://www.velodynelidar.com/>

⁶http://www.asus.com/Multimedia/Xtion_PRO_LIVE/




Fig. 2. *People detection*: The developed RefleX camera system allows to detect workers in different poses such as standing (Left) or sitting (Right).

manipulator motions by formulating a representation of the corresponding tasks as a hierarchical Stack of Tasks (SoT) which is executed by the prioritized motion control framework proposed by Kanoun et al. [13].

A. Challenges in Autonomous Grasping

In current autonomous grasping systems [14], [15], [16], grasp planning and manipulator motion planning are usually seen as independent sub-problems. A database storing object models together with sets of pre-computed grasps is used to find suitable gripper poses and joint configurations. In the online stage, sampling based planners [17] attempt to generate valid trajectories for the pre-planned grasps, which are executed in a feasible-first manner [14]. During the execution phase, such approaches necessitate many futile motion planning attempts, which often incurs significant time delays since sampling based planners suffer from the curse of dimensionality. Also, while being able to solve complicated planning problems if given enough time, these planners do not scale well to geometrically simple scenarios and they are ill suited to incorporate contact events with the environment which is a prerequisite for any grasping/manipulation application.

Using control to exploit manipulator redundancy has long been in the focus of research [18], [19]. In this line of thought, opposed to define a set of discrete grasp poses, Gienger et al. [20], [21] learn a manifold of feasible grasps for an object and use an attractor-based control scheme to reach the corresponding volume in the task space. Another approach to address motion planning and motion control simultaneously, is to formulate a hierarchical stack of equality tasks which are represented as task functions [22]. Lower-ranked tasks are then solved in the null-space of tasks with higher priority [18], [19]. A major step forward in this direction was achieved by Kanoun et al. [13], who extended the approach to include inequality tasks which significantly increases the expressiveness of the method. Recent advancements include the extension to dynamic control [23] and an efficient solution algorithm for the underlying optimization problem [24].

Algorithm 1: Object detection algorithm

```

input : Pointcloud  $\mathcal{P}$ , transform from sensor frame to
        world frame  ${}^w_sT(q)$ 
output: Target clusters  $\{\mathcal{T}\}$ 
Transform  $\mathcal{P}$  into world coordinates using  ${}^w_sT(q)$ ;
while target plane not found do
    Extract pallet plane normal  $\mathbf{n}$ , offset  $d$  from  $\mathcal{P}$ 
    using RANSAC;
    if  $(\mathbf{n}^T \mathbf{e}_z < \cos \alpha) \wedge (abs(d - h_p) < \epsilon_p)$  then
        | target plane found;
    end
    else
        |  $\mathcal{P} \leftarrow$  remove inliers;
        | if  $|\mathcal{P}|$  less than 20% of the original then
            | report failure;
            | return;
        | end
    end
end
 $\mathcal{B} \leftarrow$  oriented bounding box of plane inliers with
 $z$ -component in  $(d, d + d_{eps})$ ;
 $\{\mathcal{E}\} \leftarrow$  euclidean cluster points in  $\mathcal{B}$ ;
for each cluster  $\mathcal{E}_i$  do
    Fit a ML Gaussian to  $\mathcal{E}_i$  as  $\mu_i, \Sigma_i$ ;
    Eigen decomposition  $\Sigma_i = Q\Lambda Q^{-1}$ ;
     $\mathbf{v}_1 \leftarrow$  eigenvector associated to the largest
    eigenvalue  $\lambda_1$ ;
    if  $(\mathbf{v}_1^T \mathbf{e}_z < \cos \epsilon_\alpha) \wedge (\lambda_2 < r_{max})$  then
        | append cluster  $\mathcal{E}_i$  to the solution set  $\{\mathcal{T}\} \leftarrow \mathcal{E}_i$ ;
    end
end

```

B. Object Detection

A typical commissioning scenario in a warehouse environment generally involves as a first step the identification and localization of an object targeted for picking. There are however many application scenario specific factors to consider when designing an object detection system for this task. For example: a single pallet may hold a homogeneous or a heterogeneous set of objects; objects may be stored in boxes and be equipped with barcodes; objects may be stored on shelves, pallets, or simply be piled up. In this work, we assume a simple scenario where objects of the same type are stored on a single pallet. As shown in Fig. 1, the gripper of the APPLE system is equipped with a Structure IO structured light depth sensor⁷. We use the 3D point clouds generated by this sensor as an input for a simple object detection module which identifies clusters of points matching our expected object models. The object detection module was implemented using the Point Cloud Library (PCL) [25] following the simple procedure summarized in Algorithm 1.

The basic idea of Algorithm 1 is to first detect the location

⁷<http://structure.io/>




Fig. 3. *Grasp interval*: The shaded cyan regions illustrate the side grasp interval constraints for a cylindrical object. For a successful grasp, the palm frame origin o needs to lie inside the depicted cylindrical shell which is aligned with object axis a . The cylinder's height is limited by two planes which are normal to a . Additionally, the gripper's vertical axis (z) is constrained to lie in a cone whose axis \hat{a} is parallel to the object axis a . Furthermore, the gripper's approach axis (x) has to lie inside a cone centered on the normal which connects axis a and point o .

of the pallet plane, then to extract the points belonging to objects on the pallet, cluster them and subsequently analyze the obtained clusters. The parameters in this algorithm are the expected height h_p of the pallet above floor level, a tolerance ϵ_p on the pallet height, a tolerance ϵ_α on the angle that the pallet plane normal makes with the vertical axis z and the maximum radius r_{max} of a graspable cylinder. Each of the detected clusters is analyzed using PCA, checking that the dominant point distribution is along the z axis and that the x/y point spread is within the graspable objects limit. Identified target clusters are then passed on to the grasp planning module described below.

C. Grasp Representation and Planning

Traditional data-driven grasp synthesis [26] relies on precise knowledge of the target object geometry and frictional properties. Since this information is often not available in practice, grasps planned in simulation and under these assumptions might fail. This holds especially true, when considering an underactuated grasping device as we do in the presented work. In this case, the gripper's joint configuration depends on the interaction with the environment and is difficult or impossible to determine at planning time. Instead, we represent grasps as predefined pose intervals associated with primitive object geometries as exemplary shown in Fig. 3 for side grasps on cylindrical objects. Note that an analogue interval can also be formed to enable top grasps. These intervals bound the grasping device's position and orientation, but do not fully constrain its pose. In this work, for




Fig. 4. *Truncated grasp interval*: During the online stage, the corresponding grasp interval shown in Fig. 3 needs to be truncated (*i. e.*, parameters for r_1 , r_2 , c , h and φ need to be determined) to accommodate the specific target object dimensions and to account for the fact that some regions of the grasp interval might not be feasible due to obstruction by the environment.

simplicity, we limit ourselves to the illustrated grasp interval defined for cylindrical shapes. Corresponding intervals can be defined for other shape categories such as spheres and parallelepipeds as well. Opposed to the similarly defined task intervals in [20], [21] which are learned in simulation, we deliberately design grasp intervals to incorporate prior knowledge about robust grasp poses in our representation. In a concept originating from observations of human grasping behavior, it has been shown that the grasping device should be roughly aligned with the target object's principal components to achieve robust grasps [27]. This property is achieved by the cone constraints for the case depicted in Fig. 3.

Currently, the parameters of the grasp intervals such as the distance range between gripper and object have to be evaluated experimentally for each primitive shape category. To ease this non-trivial requirement, in the presented work we rely on a gripper which offers a low pre-grasp pose sensitivity combined with a compliant and robust grasp execution routine as discussed in Section III-E.

During operation, after the target object pose is detected as the mean and covariance of the cluster obtained from Algorithm 1, the grasp interval needs to be adapted to the specific scene and target object dimensions as illustrated in Fig. 4. For the evaluation presented in Section IV we pre-defined the corresponding parameters and gripper pre-grasp joint configuration, an appropriate programmatic approach is left to future work. One promising strategy currently being explored is to pre-compute a gripper collision map and use it for determining the valid grasping configurations around the target object. Once invalid grasping configurations have been pruned by using the scene model, the remaining config-

urations can be clustered into continuous regions and used to determine the constraints in an automatic and collision-aware manner.

D. Manipulator Motion Generation

To leverage the freedom in pre-grasp position and orientation gained from the previously described grasp representation, we use a control-based approach for reactive, on-the-fly motion generation. More specific, the method of choice is the one developed by Kanoun et al. [13] which allows to formulate the grasp interval, as well as additional desiderata such as joint limit avoidance, in form of task functions which subsequently are utilized to form a hierarchical SoT. In this paradigm, the SoT is then used to compute prioritized controls during movement execution. Therefore, motions are generated instantaneously without the planning delays occurring in sense-plan-act architectures. In the following, we briefly revise the concept. For an in-depth discussion the reader is referred to the original work presented in [13].

We lean on the notation in [24] and define the manipulator joint configuration by the vector \mathbf{q} and the control inputs as corresponding joint velocities $\dot{\mathbf{q}}$. A task function is any derivable function of \mathbf{e} . To give an example, a task with the purpose of bringing an end-effector point $\mathbf{p}(\mathbf{q})$ onto a plane described by unit normal \mathbf{n} and offset d can be transcribed by the task function $\mathbf{e} = \mathbf{n}^T \mathbf{p}(\mathbf{q}) - d$, which formulates the projection residual between the plane and $\mathbf{p}(\mathbf{q})$. The task evolution is given by $\mathbf{J}\dot{\mathbf{q}} = \dot{\mathbf{e}}$ with task jacobian $\mathbf{J} = \frac{\partial \mathbf{e}}{\partial \mathbf{q}}$.

The goal is to compute joint velocities such that the task evolution follows a desired reference profile $\dot{\mathbf{e}}^*$ (in this work chosen as exponential decay $\dot{\mathbf{e}}^* = -\lambda \mathbf{e}$, with $\lambda \in \mathbb{R}_+$). For a single equality task, this necessitates to solve the following least squares Quadratic Program (QP)

$$\dot{\mathbf{q}}^* \in \arg \min_{\dot{\mathbf{q}}} \|\mathbf{J}\dot{\mathbf{q}} - \dot{\mathbf{e}}^*\|. \quad (1)$$

In order to allow for inequality tasks, we henceforth use a general task formulation with upper bounds

$$\mathbf{J}\dot{\mathbf{q}} \leq \dot{\mathbf{e}}^*. \quad (2)$$

As stated in [24], this allows to transcribe lower bounds $\mathbf{J}\dot{\mathbf{q}} \geq \dot{\mathbf{e}}^*$, double bounds $\underline{\dot{\mathbf{e}}}^* \leq \mathbf{J}\dot{\mathbf{q}} \leq \dot{\mathbf{e}}^*$ and equalities $\mathbf{J}\dot{\mathbf{q}} = \dot{\mathbf{e}}^*$ by reformulating the task respectively as $-\mathbf{J}\dot{\mathbf{q}} \leq -\dot{\mathbf{e}}^*$, $\begin{bmatrix} -\mathbf{J} \\ \mathbf{J} \end{bmatrix} \dot{\mathbf{q}} \leq \begin{bmatrix} -\dot{\mathbf{e}}^* \\ \dot{\mathbf{e}}^* \end{bmatrix}$ and $\begin{bmatrix} -\mathbf{J} \\ \mathbf{J} \end{bmatrix} \dot{\mathbf{q}} \leq \begin{bmatrix} -\dot{\mathbf{e}}^* \\ \dot{\mathbf{e}}^* \end{bmatrix}$.

If the constraint in (2) is infeasible, a least squares solution for \mathbf{q}^* as in (1) can be found by introducing the slack variable \mathbf{w} in the decision variables

$$\begin{aligned} & \min_{\dot{\mathbf{q}}, \mathbf{w}} \|\mathbf{w}\| \\ & \text{subject to } \mathbf{J}\dot{\mathbf{q}} \geq \mathbf{e}^* + \mathbf{w}. \end{aligned} \quad (3)$$

To form a hierarchical SoT with $p = 1, \dots, P$ priority levels, we stack all task jacobians in (2) with the same assigned priority in a matrix \mathbf{A}_p , and all corresponding reference velocities in a vector \mathbf{b}_p to form one constraint of the form $\mathbf{A}_p \dot{\mathbf{q}} \leq \mathbf{b}_p$ for each hierarchy level. The aim is to sequentially satisfy a constraint at best in the least square sense while




Fig. 5. *Pull-in grasping strategy*: Depicted is a sequence of intermediate grasp states where the belts of the gripper are used to pull the object towards its palm which results in a transition from a fingertip to an enveloping grasp.

solving for the subsequent constraint of lower priority in the null-space of the previous constraint, such that the previous solution is left unchanged. Therefore, the following QP, where the previous slack variable solutions \mathbf{w}_i^* are frozen between iterations, needs to be solved for $p = 1, \dots, P$

$$\begin{aligned} & \min_{\dot{\mathbf{q}}, \mathbf{w}_p} \|\mathbf{w}_p\| \\ & \text{subject to } \mathbf{A}_i \dot{\mathbf{q}} \leq \mathbf{b}_i + \mathbf{w}_i^*, \quad i = 1, \dots, p-1 \\ & \mathbf{A}_p \dot{\mathbf{q}} \leq \mathbf{b}_p + \mathbf{w}_p. \end{aligned} \quad (4)$$

The final control vector $\dot{\mathbf{q}}^*$ is then obtained from the P th solution of (4). Opposed to sampling based planners, the chosen control scheme allows to incorporate qualitative requirements (*e.g.*, the desired gripper alignment) during motion generation and allows for redundancy exploitation via appropriate task function formulations.

Manipulator obstacle avoidance is also achieved on a control-level, by formulating tasks which maintain minimum distances between simple geometric primitives such as spheres, planes, points and capsules. We argue that for the considered application strict collision avoidance is neither necessary nor desired, since picking and manipulation inherently necessitates contact events between the robot and the environment. Also, in real-world applications where knowledge about the environment is available only in form of noisy sensor data, it might not be possible to avoid contacting the environment without being overly conservative. To this end, the APPLE platform relies on the compliant low-level control schemes and contact detection abilities of the employed manipulator.

E. Robust Grasp Execution

For this component, we exploit the capabilities of the utilized gripper [3], namely underactuation and conveyor belts on the finger pads in order to achieve robust grasping behavior. Each of the grippers two fingers has a planar manipulator structure with two rotary joints and active surfaces which are implemented by conveyor belts on the inside of the two phalanges. The mechanical structure is underactuated and comprises only one actuated Degree of Freedom (DoF) for opening and closing and two DoF per finger for the belt movements. If, during grasping, the proximal phalanges are blocked by an object, the grippers distal phalanges continue to “wrap around” and envelope it in a firm grasp. The experiments reported in [16] showed, that in cluttered scenes fingertip grasps are more likely to be feasible than robust enveloping grasps, because the latter necessitate large opening angles resulting in bulky gripper silhouettes for




Fig. 6. *Grasp Execution Control*: (Left) as the gripper starts closing, the current through the opening motor is monitored; (Middle) when contact is made, the actuated belts are switched on to pull in the object; (Right) the controller strives to maintain a given current setpoint to enable compliant in-hand manipulation behavior;

which there might no collision free approach trajectory exist. Therefore, we employ the “pull-in” strategy which is illustrated in Fig. 5. As demonstrated previously [28], this strategy is especially effective in achieving stable grasps, starting from a relatively wide range of initial gripper poses with respect to the target object.

The grasping controller was implemented by means of low-level current control of the gripper’s actuators. Since the gripper drives have a low transmission ratio and are easily back-driveable, current control enables a simple compliant behavior because the current absorption increases with increasing effort on the output. Thus, in the present case, the motor current is proportional to the resulting grasping force. The developed grasping strategy is implemented in three steps as illustrated in Fig. 6. First, a relatively low current setpoint is given to the open/close motor controller and the fingers start closing. Once the fingers come into contact with the object and the target current is reached, the motion stops which is detected by monitoring the corresponding motor encoder. Now the second phase of the grasping process is triggered and the belts are actuated to pull in the target object. We then monitor the encoders on the belts and on the phalange joints. If the belts block and the phalanges have wrapped around the object, an enveloping grasp is achieved and we transition to the final stage. Here, a higher current setpoint is given to the open/close motor controller to ensure a firm grasp. The main parameters to this routine (the current thresholds for contact and final enveloping grasp) depend on the target object properties – namely friction coefficient and mass. In this work, the parameters were experimentally tuned and good values were found for a set of common household objects.

IV. EVALUATION

For an early evaluation of the system, we set up a simplified commissioning scenario as depicted in Fig. 7. To this end, the previously described software components were implemented in the Robot Operating System (ROS) [29] framework. For manipulator motion control (see Section III-D) we use the task function formulations in [30] for joint limit avoidance, obstacle avoidance and to formulate the

inequality constraints illustrated in Fig. 4 in order for the manipulator to reach the grasp interval. An off-the-shelf solver [31] is used to carry out the optimizations for the motion control according to (4). The run time for the whole procedure is approximately 4 minutes of which 2 minutes are spent on object detection and grasping (see the attached video submission).

V. DISCUSSION AND OUTLOOK

In this paper, we presented a research platform for autonomous picking and palletizing (APPLE) for fully autonomous commissioning tasks in intralogistics settings. The APPLE robot comprises a nonholonomic mobile base with the ability to autonomously detect and pick standard EUR half-pallets from designated loading areas. Also incorporated is a camera system for the detection and avoidance of human workers wearing reflective vests. The manipulation system for loading/unloading unstructured goods from pallets operates on a novel redundant grasp representation as intervals in task space, which allows to incorporate empirical knowledge. We leverage the obtained redundancy by generating reactive manipulator motions on the fly, by using a prioritized control approach which allows to formulate the target object picking as a stack of hierarchical tasks [13]. Furthermore, we provide an early experimental evaluation of the APPLE system by means of a simplified commissioning task (see the video attachment to this article).

The presented work is limited to the picking of objects with cylindrical shapes, future work will aim at extending our grasp interval representation to other primitive shape types. Also, we work on a computational method to parametrize the interval for a current scene which, right now, is done empirically.

ACKNOWLEDGMENTS

The authors would like to thank Per Sporrang, Bo-Lennart Silfverdal, Bengt Åsberg and Joakim Larsson at Örebro University for their engineering support.



Fig. 7. *Evaluation setup*: (Left) the robot picks up an empty pallet in a designated zone; (Middle) the robot navigates to a loading zone where a can is detected and picked up; (Right) the loaded pallet is transported to a drop-off location.

REFERENCES

- [1] W. Echelmeyer, A. Kirchheim, and E. Wellbrock, "Robotics-logistics: Challenges for automation of logistic processes," in *Proc. of the IEEE International Conference on Automation and Logistics*, 2008, pp. 2099–2103.
- [2] P. R. Wurman, R. D'Andrea, and M. Mountz, "Coordinating hundreds of cooperative, autonomous vehicles in warehouses," *AI magazine*, vol. 29, no. 1, pp. 9–20, 2008.
- [3] V. Tincani, M. G. Catalano, E. Farnioli, M. Garabini, G. Grioli, G. Fantoni, and A. Bicchi, "Velvet fingers: A dexterous gripper with active surfaces," in *Proc. of the IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2012, pp. 1257–1263.
- [4] T. Hellstrom and O. Ringdahl, "Follow the past: a path-tracking algorithm for autonomous vehicles," *Int. Journal of Vehicle Autonomous Systems*, vol. 4, no. 2, pp. 216–224, 2006.
- [5] J. Marshall, T. Barfoot, and J. Larsson, "Autonomous underground tramming for center-articulated vehicles," *Journal of Field Robotics*, vol. 25, no. 6-7, pp. 400–421, 2008.
- [6] K. Hyypä, "Method of navigating an automated guided vehicle," Mar. 7 1989, uS Patent 4,811,228.
- [7] M. Cirillo, T. Uras, S. Koenig, H. Andreasson, and F. Pecora, "Integrated motion planning and coordination for industrial vehicles," in *Proc. of the International Conference on Automated Planning and Scheduling*, 2014.
- [8] H. Andreasson, J. Saarinen, M. Cirillo, T. Stoyanov, and A. J. Lilienthal, "Fast, continuous state path smoothing to improve navigation accuracy," in *Proc. of the IEEE International Conference on Robotics and Automation*, 2015, to appear.
- [9] T. Stoyanov, J. P. Saarinen, H. Andreasson, and A. J. Lilienthal, "Normal distributions transform occupancy map fusion: Simultaneous mapping and tracking in large scale dynamic environments," in *Proc. of the IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2013, pp. 4702–4708.
- [10] R. Valencia, J. Saarinen, H. Andreasson, J. Vallve, J. Andrade-Cetto, and A. Lilienthal, "Localization in highly dynamic environments using dual-timescale ndt-mcl," in *Proc. of the IEEE International Conference on Robotics and Automation*, May 2014, pp. 3956–3962.
- [11] D. R. Canelhas, T. Stoyanov, and A. J. Lilienthal, "Sdf tracker: A parallel algorithm for on-line pose estimation and scene reconstruction from depth images," in *IROS*. IEEE, 2013, pp. 3671–3676.
- [12] R. Mosberger, H. Andreasson, and A. J. Lilienthal, "A customized vision system for tracking humans wearing reflective safety clothing from industrial vehicles and machinery," *Sensors*, vol. 14, no. 10, pp. 17952–17980, 2014.
- [13] O. Kanoun, F. Lamiraux, and P.-B. Wieber, "Kinematic control of redundant manipulators: Generalizing the task-priority framework to inequality task," *IEEE Transactions on Robotics*, vol. 27, no. 4, pp. 785–792, 2011.
- [14] D. Berenson, R. Diankov, K. Nishiwaki, S. Kagami, and J. Kuffner, "Grasp planning in complex scenes," in *Proc. IEEE/RAS International Conference on Humanoid Robots*, 2007, pp. 42–48.
- [15] S. Srinivasa, D. Ferguson, C. Helfrich, D. Berenson, A. Collet, R. Diankov, G. Gallagher, G. Hollinger, J. Kuffner, and M. VandeWeghe, "Herb: A home exploring robotic butler," *Autonomous Robots*, vol. 28, no. 1, pp. 5–20, 2010.
- [16] R. Krug, T. Stoyanov, M. Bonilla, V. Tincani, N. Vaskevicius, G. Fantoni, A. Birk, A. J. Lilienthal, and A. Bicchi, "Velvet fingers: Grasp planning and execution for an underactuated gripper with active surfaces," in *Proc. of the IEEE International Conference on Robotics and Automation*, 2014, pp. 3669–3675.
- [17] S. M. LaValle, *Planning Algorithms*. Cambridge University Press, 2006.
- [18] B. Siciliano and J.-J. Slotine, "A general framework for managing multiple tasks in highly redundant robotic systems," in *Proc. of the International Conference on Advanced Robotics*. IEEE, 1991, pp. 1211–1216.
- [19] L. Sentis, J. Park, and O. Khatib, "Compliant control of multicontact and center-of-mass behaviors in humanoid robots," *IEEE Transactions on Robotics*, vol. 26, no. 3, pp. 483–501, 2010.
- [20] M. Gienger, M. Toussaint, and C. Goerick, "Task maps in humanoid robot manipulation," in *Proc. of the IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2008, pp. 2758–2764.
- [21] M. Gienger, M. Toussaint, N. Jetchev, A. Bendig, and C. Goerick, "Optimization of fluent approach and grasp motions," in *Proc. IEEE/RAS International Conference on Humanoid Robots*, 2008, pp. 111–117.
- [22] C. Samson, B. Espiau, and M. L. Borge, *Robot control: the task function approach*. Oxford University Press, 1991.
- [23] L. Saab, O. E. Ramos, F. Keith, N. Mansard, P. Soueres, and J. Fourquet, "Dynamic whole-body motion generation under rigid contacts and other unilateral constraints," *IEEE Transactions on Robotics*, vol. 29, no. 2, pp. 346–362, 2013.
- [24] A. Escande, N. Mansard, and P.-B. Wieber, "Hierarchical quadratic programming: Fast online humanoid-robot motion generation," *International Journal of Robotics Research*, 2014.
- [25] R. B. Rusu and S. Cousins, "3D is here: Point Cloud Library (PCL)," in *Proc. of the IEEE International Conference on Robotics and Automation*, 2011, pp. 1–4.
- [26] J. Bohg, A. Morales, T. Asfour, and D. Kragic, "Data-driven grasp synthesis: a survey," *IEEE Transactions on Robotics*, vol. 30, no. 2, pp. 289–309, 2014.
- [27] R. Balasubramanian, L. Xu, P. Brook, J. Smith, and Y. Matsuoaka, "Physical human interactive guidance: Identifying grasping principles from human-planned grasps," *IEEE Transactions on Robotics*, vol. 28, no. 4, pp. 899–910, 2012.
- [28] R. Krug, T. Stoyanov, M. Bonilla, V. Tincani, N. Vaskevicius, G. Fantoni, A. Birk, A. J. Lilienthal, and A. Bicchi, "Improving grasp robustness via in-hand manipulation with active surfaces," in *IEEE International Conference on Robotics and Automation - Workshop on Autonomous Grasping and Manipulation: An Open Challenge*, 2014.
- [29] M. Quigley, K. Conley, B. Gerkey, J. Faust, T. Foote, J. Leibs, R. Wheeler, and A. Y. Ng, "ROS: an open-source robot operating system," in *IEEE International Conference on Robotics and Automation - Workshop on open source software*, vol. 3, no. 3.2, 2009, p. 5.
- [30] O. Kanoun, "Contribution à la planification de mouvement pour robots humanoïdes," Ph.D. dissertation, l'Université Toulouse III - Paul Sabatier, 2009.
- [31] I. Gurobi Optimization, "Gurobi optimizer reference manual," 2015. [Online]. Available: <http://www.gurobi.com>